

PS7_Solutions

February 27, 2019

0.0.1 Problem Set 7 Keertana V. Chidambaram

Problem 1

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, LeaveOneOut, KFold, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, mean_squared_error
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
from matplotlib.ticker import MultipleLocator
import statsmodels.api as sm
import scipy.interpolate as si
```

```
In [2]: data = pd.read_csv('data/strongdrink.txt')
```

```
In [3]: data.head()
```

```
Out[3]:
```

	cultivar	alco	malic	ash	alk	magn	tot_phen	flav	nonfl_phen	\
0	1	14.23	1.71	2.43	15.6	127	2.80	3.06	0.28	
1	1	13.20	1.78	2.14	11.2	100	2.65	2.76	0.26	
2	1	13.16	2.36	2.67	18.6	101	2.80	3.24	0.30	
3	1	14.37	1.95	2.50	16.8	113	3.85	3.49	0.24	
4	1	13.24	2.59	2.87	21.0	118	2.80	2.69	0.39	

	proanth	color_int	hue	OD280rat	proline
0	2.29	5.64	1.04	3.92	1065
1	1.28	4.38	1.05	3.40	1050
2	2.81	5.68	1.03	3.17	1185
3	2.18	7.80	0.86	3.45	1480
4	1.82	4.32	1.04	2.93	735

```
In [4]: #Solution 1.a.
y = data['cultivar']
x_vars = ['alco', 'malic', 'tot_phen', 'color_int']
X = data[x_vars]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 42)
```

```
In [5]: LogReg = LogisticRegression(solver='newton-cg',multi_class='multinomial').fit(X_train,
index = pd.DataFrame(['beta0', 'beta1', 'beta2', 'beta3', 'beta4']))
j1 = pd.DataFrame([LogReg.intercept_[0]] + list(LogReg.coef_[0]))
j2 = pd.DataFrame([LogReg.intercept_[1]] + list(LogReg.coef_[1]))
result = pd.concat([index, j1, j2], axis=1)
result.columns = ['coefficient', 'j = 1', 'j = 2']
result
```

```
Out[5]:
```

	coefficient	j = 1	j = 2
0	beta0	-24.010989	22.802446
1	beta1	1.700403	-1.468044
2	beta2	-0.265605	-0.333053
3	beta3	1.223894	0.664012
4	beta4	0.022756	-0.922712

```
In [6]: y_pred=LogReg.predict(X_test)
target_names = ['j=1', 'j=2', 'j=3']
print(classification_report(y_test, y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
j=1	0.87	1.00	0.93	13
j=2	1.00	0.90	0.95	21
j=3	1.00	1.00	1.00	10
micro avg	0.95	0.95	0.95	44
macro avg	0.96	0.97	0.96	44
weighted avg	0.96	0.95	0.96	44

```
In [7]: error_rates = pd.DataFrame([1 - 0.87, 0, 0])
error_rates.index = target_names
error_rates.columns = ['error rate']
error_rates
```

```
Out[7]:
```

	error rate
j=1	0.13
j=2	0.00
j=3	0.00

j=2,3 are the categories with the lowest error rates. But j=3 has the best f1 score. So j=3 has best prediction according to this model.

```
In [8]: data['cultivar'].value_counts()
```

```
Out[8]:
```

2	71
1	59
3	46

Name: cultivar, dtype: int64

j=2 has the most number of observations but it is not the best predicted group in this model.

```
In [9]: print('Test set MSE = ', mean_squared_error(y_test, y_pred))
```

```
Test set MSE = 0.045454545454545456
```

```
In [10]: #Solution 1.b.
```

```
In [11]: Xvars = X.values
```

```
         yvars = y.values
```

```
         N_loo = Xvars.shape[0]
```

```
         loo = LeaveOneOut()
```

```
         loo.get_n_splits(Xvars)
```

```
         MSE_vec = np.zeros(N_loo)
```

```
         y_pred = np.zeros(N_loo)
```

```
         for train_index, test_index in loo.split(Xvars):
```

```
             X_train, X_test = Xvars[train_index], Xvars[test_index]
```

```
             y_train, y_test = yvars[train_index], yvars[test_index]
```

```
             LogReg = LogisticRegression(solver='newton-cg', multi_class='multinomial').fit(X_train, y_train)
```

```
             y_pred[test_index] = LogReg.predict(X_test)
```

```
             MSE_vec[test_index] = y_test != y_pred[test_index]
```

```
         MSE_loo = MSE_vec.mean()
```

```
         print('test estimate MSE loocv=', MSE_loo)
```

```
test estimate MSE loocv= 0.07954545454545454
```

```
In [12]: print(classification_report(yvars, y_pred))
```

	precision	recall	f1-score	support
1	0.90	0.93	0.92	59
2	0.91	0.90	0.91	71
3	0.96	0.93	0.95	46
micro avg	0.92	0.92	0.92	176
macro avg	0.92	0.92	0.92	176
weighted avg	0.92	0.92	0.92	176

```
In [13]: error_rates = pd.DataFrame([1 - 0.90, 1 - 0.91, 1 - 0.96])
```

```
         error_rates.index = target_names
```

```
         error_rates.columns = ['error rate']
```

```
         error_rates
```

```
Out[13]:      error rate
      j=1      0.10
      j=2      0.09
      j=3      0.04
```

Clearly, the error rates for all 3 categories have increased as compared to part a.

```
In [14]: #Solution 1.c.
```

```
In [15]: Xvars = X.values
        yvars = y.values
        k = 4
        kf = KFold(n_splits=4, shuffle=True, random_state=10)
        kf.get_n_splits(Xvars)

        MSE_vec_kf = np.zeros(k)
        y_pred = np.zeros(len(yvars))

        k_ind = int(0)
        for train_index, test_index in kf.split(Xvars):
            X_train, X_test = Xvars[train_index], Xvars[test_index]
            y_train, y_test = yvars[train_index], yvars[test_index]
            LogReg = LogisticRegression(solver='newton-cg', multi_class='multinomial').fit(X_train, y_train)

            y_pred[test_index] = LogReg.predict(X_test)

            err = y_pred[test_index] != y_test
            MSE_vec_kf[k_ind] = err.mean()
            k_ind += 1

        MSE_kf = MSE_vec_kf.mean()
        print('test estimate MSE k-fold (k=4) =', MSE_kf)
```

```
test estimate MSE k-fold (k=4) = 0.09090909090909091
```

```
In [16]: print(classification_report(yvars, y_pred))
```

	precision	recall	f1-score	support
1	0.87	0.93	0.90	59
2	0.91	0.87	0.89	71
3	0.96	0.93	0.95	46
micro avg	0.91	0.91	0.91	176
macro avg	0.91	0.91	0.91	176
weighted avg	0.91	0.91	0.91	176

```
In [17]: error_rates = pd.DataFrame([1 - 0.87, 1 - 0.91, 1 - 0.96])
        error_rates.index = target_names
        error_rates.columns = ['error rate']
        error_rates
```

```
Out[17]:      error rate
j=1      0.13
j=2      0.09
j=3      0.04
```

The error rates are the same for $j = 2, 3$ but it has increased for $j=1$.

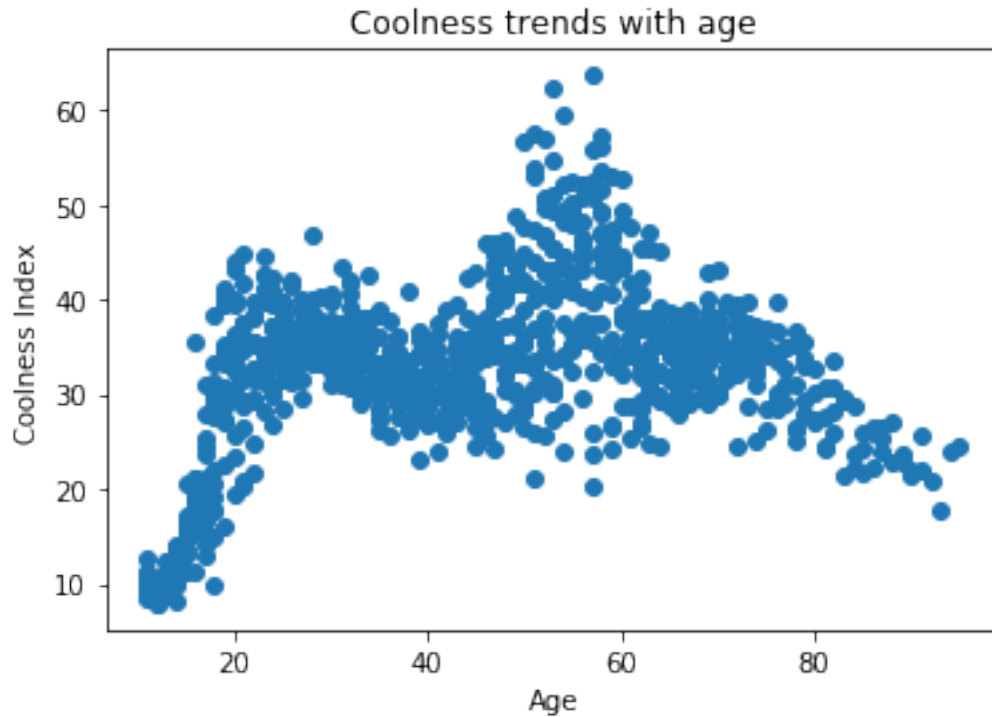
Problem 2

```
In [18]: data = pd.read_csv('data/CoolIndex.txt', names = ['age', 'coolness'])
```

```
In [19]: data.head()
```

```
Out[19]:   age  coolness
0  11.0  10.981602
1  11.0  11.364925
2  11.0  10.190227
3  11.0   9.903725
4  11.0   8.997918
```

```
In [20]: #Solution 2.a.
        X = data['age']
        y = data['coolness']
        plt.scatter(X, y)
        plt.xlabel('Age')
        plt.ylabel('Coolness Index')
        plt.title('Coolness trends with age')
        plt.show()
```




```
In [21]: #Solution 2.b.
```

```
In [22]: data['bin1'] = 1
data['bin1'] = data['bin1'].where((data.age >= 11) & (data.age < 22), 0)
data['bin2'] = 1
data['bin2'] = data['bin2'].where((data.age >= 22) & (data.age < 40), 0)
data['bin3'] = 1
data['bin3'] = data['bin3'].where((data.age >= 40) & (data.age < 59), 0)
data['bin4'] = 1
data['bin4'] = data['bin4'].where((data.age >= 59) & (data.age < 77), 0)
data['bin5'] = 1
data['bin5'] = data['bin5'].where((data.age >= 77) & (data.age <= 95), 0)
```

```
In [23]: xvars = ['bin1', 'bin2', 'bin3', 'bin4', 'bin5']
X = data[xvars]
y = data['coolness']
result = sm.OLS(y, X).fit()
```

```
In [24]: X = data['age']
y = data['coolness']
plt.scatter(X, y, label='Scatterplot')
plt.plot(X, result.predict(), color = 'm', label = "Step function", linewidth=3.0)
plt.xlabel('Age')
plt.ylabel('Coolness Index')
```

```
plt.title('Coolness trends with age')
plt.legend()
plt.show()
```



```
In [25]: result.summary()
```

```
Out[25]: <class 'statsmodels.iolib.summary.Summary'>
        """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  coolness    R-squared:                0.429
Model:                            OLS      Adj. R-squared:           0.427
Method:                 Least Squares    F-statistic:                178.7
Date:                Wed, 27 Feb 2019    Prob (F-statistic):        3.73e-114
Time:                        01:28:08    Log-Likelihood:            -3214.5
No. Observations:                  956    AIC:                        6439.
Df Residuals:                      951    BIC:                        6463.
Df Model:                            4
Covariance Type:                  nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
bin1	20.1025	0.562	35.746	0.000	18.999	21.206
bin2	34.4758	0.431	80.006	0.000	33.630	35.321

bin3	37.6351	0.424	88.814	0.000	36.804	38.467
bin4	35.2254	0.485	72.560	0.000	34.273	36.178
bin5	27.2964	0.936	29.175	0.000	25.460	29.132
=====						
Omnibus:		80.102	Durbin-Watson:		1.236	
Prob(Omnibus):		0.000	Jarque-Bera (JB):		101.718	
Skew:		0.714	Prob(JB):		8.17e-23	
Kurtosis:		3.719	Cond. No.		2.21	
=====						

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
"""
```

Coefficients for bin 1 to 5 correspond to the estimated value for beta 1 to 5 respectively. A 73 year old belongs in bin 4, thus predicted value = 35.2254, alternatively we can also show:

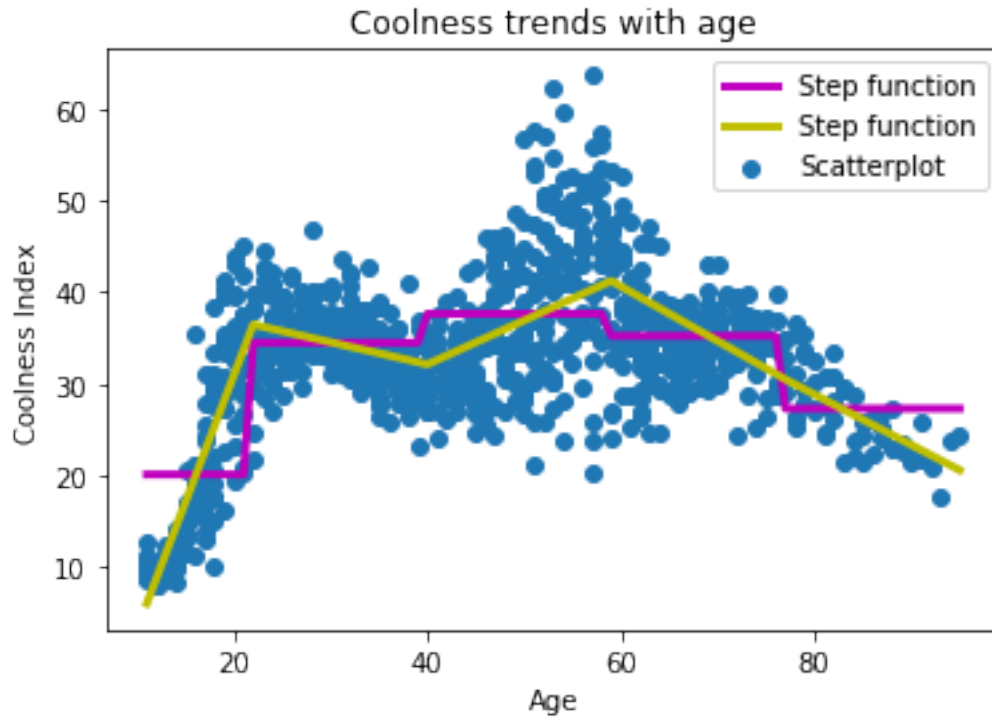
```
In [26]: result.predict([0, 0, 0, 1, 0])
```

```
Out[26]: array([35.22540004])
```

```
In [27]: #Solution 2.c.
```

```
In [28]: data2 = data.groupby('age').mean()
data2['age'] = data2.index
knots = [22, 40, 59, 77]
spl_lin_cool = si.LSQUnivariateSpline(np.array(data2['age']), np.array(data2['coolness']))
```

```
In [29]: X = data['age']
y = data['coolness']
X2 = data2['age']
plt.scatter(X, y, label='Scatterplot')
plt.plot(X, result.predict(), color = 'm', label = "Step function", linewidth=3.0)
plt.plot(X2, spl_lin_cool(X2), color = 'y', label = "Step function", linewidth=3.0)
plt.xlabel('Age')
plt.ylabel('Coolness Index')
plt.title('Coolness trends with age')
plt.legend()
plt.show()
```

Predicted coolness of a 73 year-old:


```
In [30]: spl_lin_cool(73)
```

```
Out[30]: array(32.86784862)
```

```
In [31]: #Solution 2.d.
```

```
In [32]: data2 = data.groupby('age').mean()
data2['age'] = data2.index
knots = [22, 40, 59, 77]
spl_cub_cool = si.LSQUnivariateSpline(np.array(data2['age']), np.array(data2['coolness']))
```

```
In [33]: X = data['age']
y = data['coolness']
X2 = data2['age']
plt.scatter(X, y, label='Scatterplot')
plt.plot(X, result.predict(), color = 'm', label = "Step function", linewidth=3.0)
plt.plot(X2, spl_lin_cool(X2), color = 'y', label = "Step function", linewidth=3.0)
plt.plot(X2, spl_cub_cool(X2), color = 'r', label = "Step function", linewidth=3.0)
plt.xlabel('Age')
plt.ylabel('Coolness Index')
plt.title('Coolness trends with age')
plt.legend()
plt.show()
```



Predicted coolness of a 73 year-old:

```
In [34]: spl_cub_cool(73)
```

```
Out[34]: array(32.64230107)
```

```
In [ ]:
```